Job Description: AI Senior Security Engineer (4-6 years)

Key Responsibilities

Build and support AI/ML pipelines, including data preprocessing, training workflows, fine-tuning, and inference optimization. Develop automation workflows, API integrations, agentic components, and deployment pipelines using Python, Docker/K8s, and CI/CD. Deploy and monitor AI workloads on cloud platforms. Implement strong security controls by aligning with OWASP Top 10, OWASP LLM Top 10, and internal AppSec requirements; identify and mitigate risks

Required Skills

AI & ML Engineering

- Basic foundational understanding of core ML concepts, including supervised and unsupervised learning, clustering techniques, basic feature engineering, and model evaluation fundamentals.
- Knowledge of TensorFlow, PyTorch, scikit-learn, Hugging Face, Transformers
- Familiarity with vector databases (FAISS, Chroma) and RAG pipelines
- Basic familiarity with Agentic workflows (Autogen, crew, langchain), multi-agent patterns, and MCP (Model Context Protocol) concepts is preferred.

Development & Automation

- Good Python skills for ML pipelines and automation
- Working knowledge of REST APIs, data processing scripts, and backend integration
- Working knowledge of Docker, K8s, Git, and CI/CD pipelines

AI Security

- Understanding of AI-specific risks: prompt injection, data leakage, jailbreak attempts, insecure RAG, insecure embeddings.
- Experience in LLM red-teaming and application of prompt guardrails.
- Good understanding of **OWASP Top 10** and **OWASP LLM Top 10** (AI/LLM-specific).
- Foundation in security, IAM, and secure secret handling.

Ops & Platform Skills

- Familiarity with deploying and monitoring AI workloads
- Familiarity with cloud platforms (GCP/AWS/Azure)